

# Databases

Week 1 – Introduction

Dr. Roger Young



# Course Overview

- Relational databases – Structured Query Language (SQL)
- ERD (Entity Relationship Diagram)
- Joins
- Java Database Connectivity (JDBC)
- Stored Procedures, Views, Triggers
- Normalization
- ACID v BASE
- NoSQL



# What is a Database?

- We know what data is. A piece of information!
- Information is not useful if not organized
- Database – a collection of data stored in an organized manner



# Why not store info in a file system?

- Still widely used today (e.g. for backup) but have **the following problems**
- **Data Redundancy** (Duplication of data)
  - Wasteful
  - Inconsistent
  - Loss of metadata integrity
    - Same data has different names in different files, or same name may have data in different files.



# Why not store info in a file system

- **Lengthy Development Times**

- Little opportunity to re-use previous development efforts

- **Limited Data Sharing**

- Users have little opportunity to share data outside their own

- **Excessive Program Maintenance**

- Factors above combine to create heavy maintenance load



# Advantages of a Database

- Minimal Data Redundancy (duplication)
- Improved Consistency
- Improved Data Sharing
- Increased Application Development Productivity
- Enforcement of Standards
- Better Data Accessibility/ Responsiveness
- Security, Backup/Recovery, Concurrency



# Advantages of a DBMS

- So why not use them always?
  - Expensive/complicated to set up & maintain
  - This cost & complexity must be offset by need
  - General-purpose, not suited for special-purpose tasks (e.g. text search!)



# Database management system

- Need for DBMS has exploded in the last years
  - **Corporate:** retail swipe/clickstreams, “customer relationship mgmt”, “supply chain mgmt”, “data warehouses”, etc.
  - **Scientific:** digital libraries, Human Genome project, NASA Missions, physical sensors, grid physics network





# Databases and roles

- Database administrators (DBAs) role
  - Design logical/physical schemas
  - Handle security and authorization
  - Data availability, crash recovery
  - Database tuning as needs evolve



# Database management system

- A **Database Management System (DBMS)** is a software system designed to store, manage, and facilitate access to databases.
- It is a data storage and retrieval system which permits data to be stored non-redundantly while making it appear to the user as if the data is well-integrated.



# Relational Database

- A relational database is a collection of data items organized as a set of formally described tables from which data can be accessed easily.
- A relational database is created using the relational model.
- The software used in a relational database is called a relational database management system (RDBMS).
- The relational database was first defined in June 1970 by of IBM



# Relational database management system

- A **data model** is a collection of concepts for describing data.
- RDBMS is based on the relational model of data
  - Each database has a set of named **relations** (**tables**)
  - Each relation has a set of named **attributes** (**columns**)
  - Each **tuple** (**row**) has a value for each attribute
  - Each attribute has a defined type



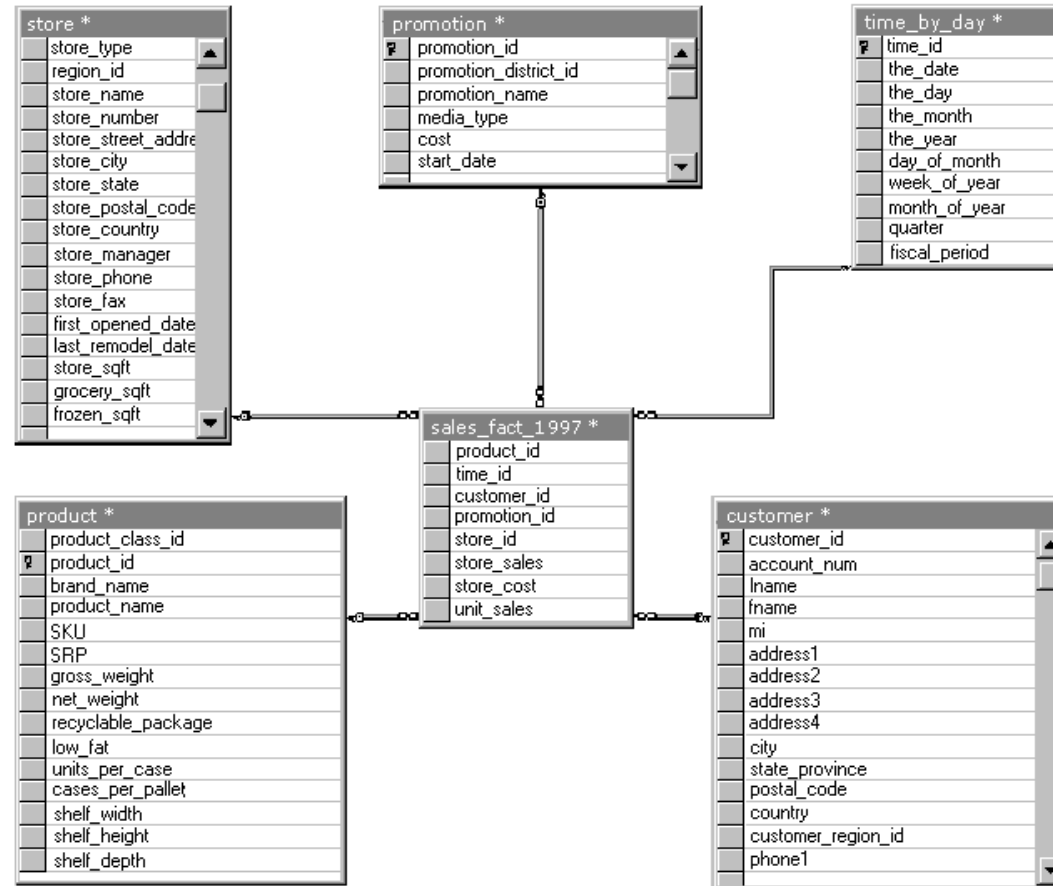
# Relational database management system

- A ***schema*** is a description of a particular collection of data, using a given data model. It describes the structure of the database.
- **Schema of RDBMS database**
  - A structural description of the relations (**tables**) in the database




# Relational database management system

- Schema Example



# Relational Databases

- A relation is defined as a set of rows that have the same attributes.
- A row usually represents an object and information about that object.
- Objects are typically physical objects or concepts.
- A relation is usually described as a table, which is organized into rows and columns.



# Definitions of Terminology

Formal relational term	Informal equivalents
relation	table
tuple	row or record
cardinality	number of rows
attribute	column or field
degree	number of columns
(unique) identifier	Primary key





# Overview of Relational Databases

- Entity
  - Object about which you want to store data
  - Different tables store data about each different entity
- Relationships
  - Links that show how different records are related



# Overview of Relational Databases

- Key fields
  - Establish relationships among records in different tables
  - Main types of key fields
    - Primary
    - Candidate
    - Foreign
    - Composite



# Primary Keys

- Column in relational database table whose value must be unique for each row
- Serves to identify individual occurrence of entity
- Every row must have a primary key
- Cannot be NULL
- NULL
  - Value is absent or unknown
  - No entry is made for that data element

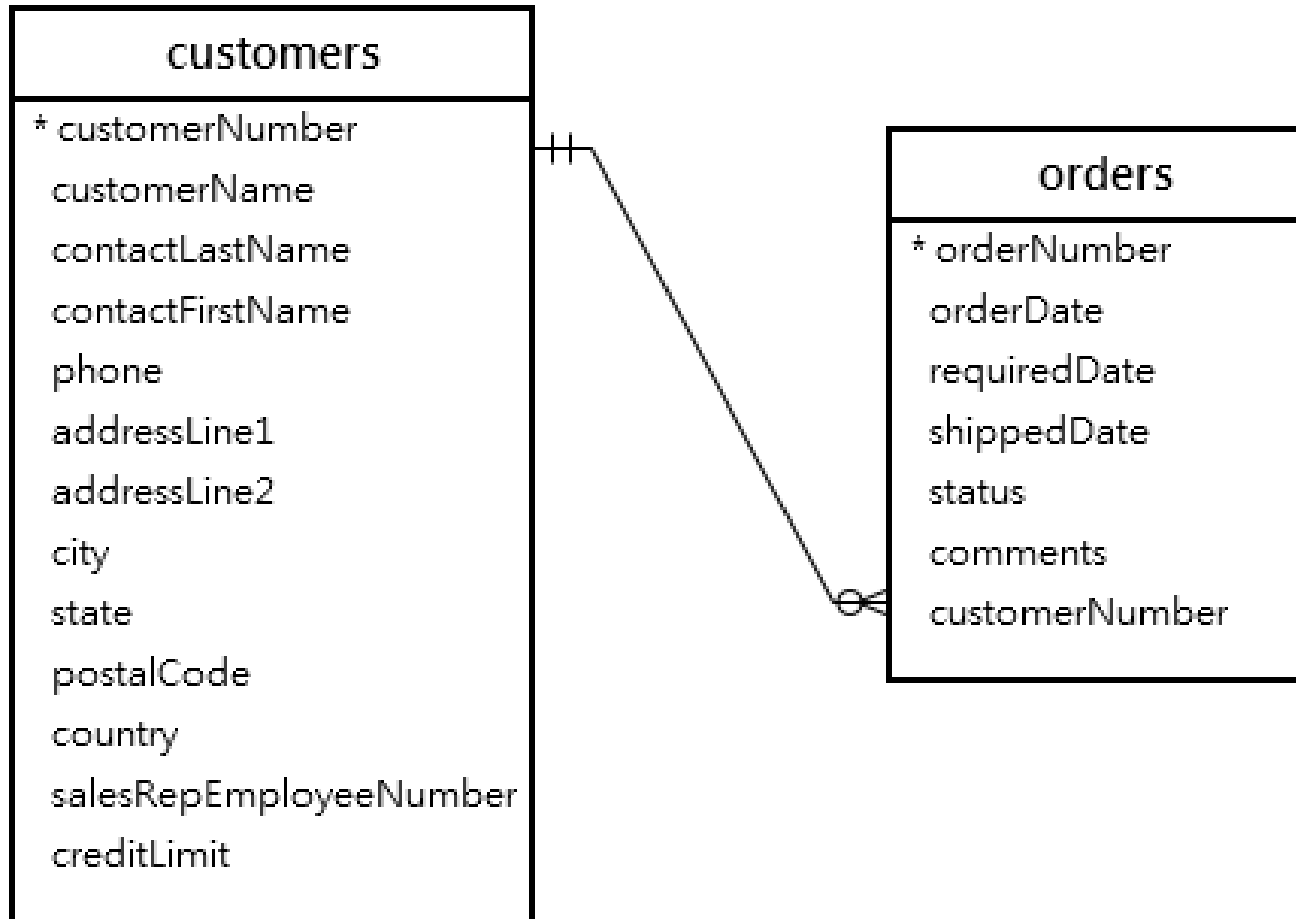


# Foreign Keys

- Column in table that is a primary key in another table
- Creates relationship between two tables
- Value must exist in table where it is the primary key



# Primary / Foreign Key Examples



# SQL Data Types

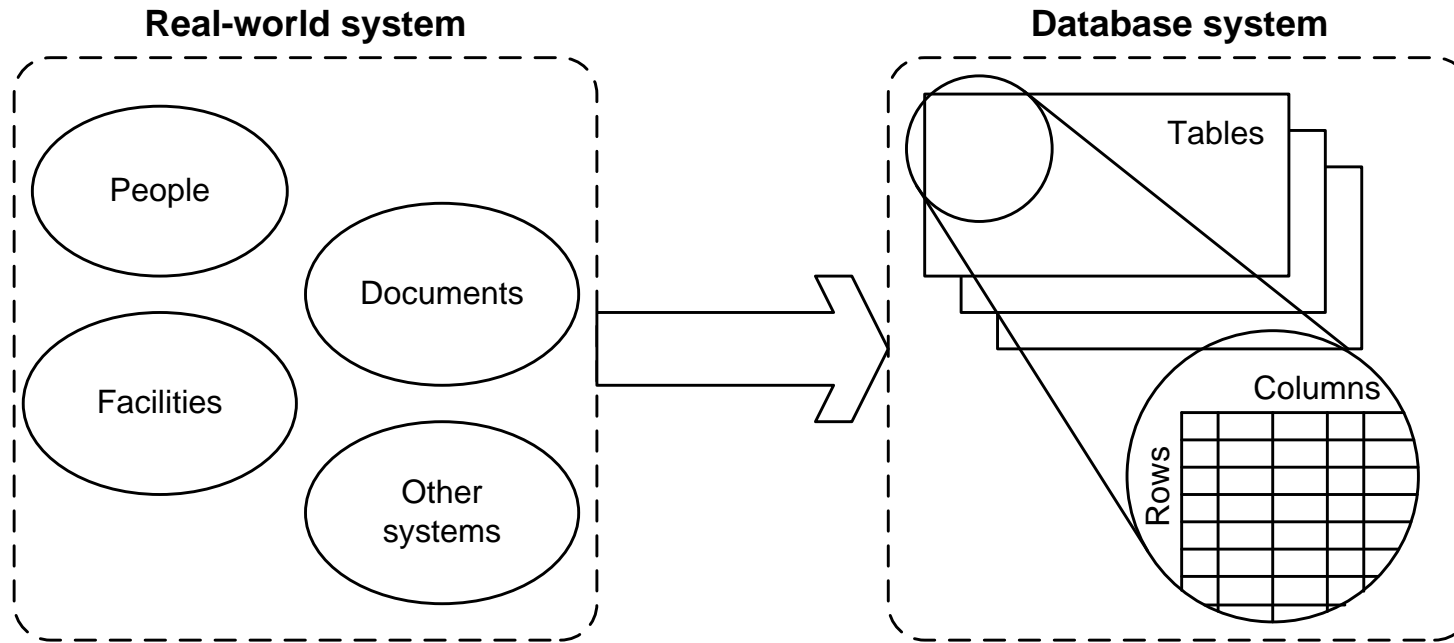
- Each Database has its own range of data types.
- In MySQL there are three main types
  - Text
  - Number
  - Date/Time



# Some MySQL Data Types

Data type	Description
CHAR(size)	Holds a <u>fixed</u> length string (can contain letters, numbers, and special characters). The fixed size is specified in parenthesis. Can store up to 255 characters
VARCHAR(size)	Holds a <u>variable</u> length string (can contain letters, numbers, and special characters). The maximum size is specified in parenthesis. Can store up to 255 characters. Note: If you put a greater value than 255 it will be converted to a TEXT type
TEXT	Holds a string with a maximum length of 65,535 characters
INT	-2147483648 to 2147483647 normal. The maximum number of digits may be specified in parenthesis
DATE()	A date. Format: YYYY-MM-DD
DATETIME()	A date and time combination. Format: YYYY-MM-DD HH:MM:SS

# A Database is modelled on a real-world system





# Database: designing the data structure

Database design	SQL statements for data definition
Identify the tables, table attributes (columns) and how tables relate to each other	CREATE DATABASE ( LIBRARY )
Identify the primary key and foreign keys, the link between the tables	CREATE TABLE For: BOOK, BORROWER, LIBRARIAN, AUTHOR, lending
Determine the data type of the attribute	
	ALTER DATABASE
	ALTER TABLE
	DROP TABLE



# SQL: data definition

<b>SHOW DATABASES;</b>	Lists all defined databases
<b>SHOW DATABASES LIKE</b> db_name;	List all defined databases that match criteria
<b>CREATE DATABASE</b> [ if not exists] db_name;	Creates database;
<b>USE</b> db_name;	All further statements are towards database
<b>CREATE TABLE</b> table_name( data_field 1 datatype, data_field2 datatype, PRIMARY KEY ( dataname));	Creates table in current database
<b>SHOW TABLES in</b> db-name;	Lists the tables



# SQL – (Structured Query Language)

- Data Definition Language (DDL)
  - Create/alter/delete tables and their attributes
- Data Manipulation Language (DML)
  - Query one or more tables
  - Insert/delete/modify tuples in tables



# Data Definition Language (DDL)

- Data Definition Language (DDL) is a vocabulary used to define data structures (database schema) in SQL. These statements are used to create, alter, or drop data structures in SQL.

1. *CREATE*

2. *DROP*

3. *ALTER*



# CREATE TABLE

- Specifies a new table by giving it a name, and specifying each of its attributes and their data types (INTEGER, FLOAT, DECIMAL(i,j), CHAR(n), VARCHAR(n))
- A constraint NOT NULL may be specified on an attribute

```
CREATE TABLE DEPARTMENT
(
    DNAME VARCHAR(10)    NOT NULL,
    DNUMBER      INTEGER NOT NULL,
    MANAGER      CHAR(9),
    MGRSTARTDATE CHAR(9) );
```



# SQL DDL Creation Syntax

```
CREATE TABLE table-name  
    (attribute-name domain,  
    attribute-name domain );
```



# SQL DDL Creation Example

```
CREATE TABLE branch  
    ( name varchar(10),  
      city varchar(20),  
      director varchar(20),  
      assets integer);
```

**branch**

name	city	director	assets



# DROP TABLE

- Used to remove a relation (table) *and its definition*
- The relation can no longer be used in queries, updates, or any other commands since its description no longer exists
- Example:

**DROP TABLE BRANCH;**





# SQL DDL Deletion Syntax

DROP TABLE *table\_name*

Examples:

DROP TABLE *branch*;

DROP TABLE *FoodCart*;



# ALTER TABLE

- Used to add an attribute to one of the tables
- The new attribute will have NULLs in all the tuples of the relation right after the command is executed; hence, the NOT NULL constraint is *not allowed* for such an attribute
- Example:

**ALTER TABLE EMPLOYEE ADD JOB VARCHAR(12);**

- The database users must still enter a value for the new attribute JOB for each EMPLOYEE tuple. This can be done using the UPDATE command.



# SQL DDL Alteration Syntax

To add an attribute:

```
ALTER TABLE table_name  
ADD Att Domain;
```

To remove an attribute:

```
ALTER TABLE table_name  
DROP Att;
```



# SQL DDL Alteration Example

ALTER TABLE *branch* **ADD** *zip* INTEGER;

<b>branch</b>	name	city	director	assets

becomes

<b>branch</b>	name	city	director	assets	zip



# SQL DDL Alteration Example

ALTER TABLE *branch* DROP *zip*;

branch	name	city	director	assets	zip



branch	name	city	director	assets



# Create your own Database & Table

- Open Workbench and create a students table with 5 columns of your choice
- To create a new database
- DROP DATABASE IF EXISTS lab1; #this will delete an existing db called test
- CREATE DATABASE IF NOT EXISTS lab1; # this will create a new db
- USE lab1; #you could have many dbs, this tells the system which one you want to work on

# Data Manipulation Language (DML)

- Data manipulation language comprises the SQL data change statements, which modify stored data but not the schema or database objects.

*INSERT INTO ... VALUES ...*

*SELECT....FROM....WHERE...*

*UPDATE ... SET ... WHERE ...*

*DELETE FROM ... WHERE ...*



# INSERT INTO (DML)

- Adds data to a table

- Syntax:

```
INSERT INTO table_name (column, ..., column)  
VALUES (value, ..., value);
```

- The *columns* are the names of columns you are putting data into, and the *values* are that data
- String data must be enclosed in single quotes
- Numbers are not quoted
- You can omit the column names if you supply a value for *every* column



# INSERT (cont.)

- In its simplest form, it is used to add one or more tuples to a relation
- Attribute values should be listed in the same order as the attributes were specified in the CREATE TABLE command

# INSERT INTO (Cont.)

- Inserting into a table
  - **Insert into employee (emp\_Name, Dept\_no, gender, salary)**  
**Values ('Sara Johns', 1, 'F', 1440);**
- Inserting a record that has some null attributes requires identifying the fields that actually get data
- When you insert a record and you have values for all attributes, there is no need to specify the attributes names.
  - **Insert into employee**  
**Values ('Suzy Alan', 10, 'F', 1200);**
- Inserting from another table
  - **INSERT INTO emp\_senior**  
**select \* from employee where age > 60;**  
**The main condition in this case, that both tables has the same attributes and ordered in the same order**

# Insert Example

- To insert a row into a table, it is necessary to have a value for each attribute, and order matters.

Example: INSERT into FoodCart  
VALUES ('02/26/08', 'pizza', 70 );

FoodCart

date	food	sold
02/25/08	pizza	350
02/26/08	hotdog	500

Becomes

date	food	sold
02/25/08	pizza	350
02/26/08	hotdog	500
02/26/08	pizza	70

# Revisiting Relational Terminology

Product

Table name

Attribute names

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Tuples or rows



# Revisiting Relational Terminology

- The *schema* of a table is the table name and its attributes:

Product(PName, Price, Category, Manufacturer)

- A *key* is an attribute whose values are unique;  
we underline a key

Product(PName, Price, Category, Manufacturer)



# SQL Query

Basic form:

```
SELECT <attributes>  
FROM   <one or more relations>  
WHERE  <conditions>
```



# Insert Values into your student table


- Insert information on 5 students into your previously created student table

# Simple SQL Query

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT *  
FROM Product  
WHERE category='Gadgets'
```



"selection"

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks






# Simple SQL Query

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT PName, Price, Manufacturer
FROM   Product
WHERE  Price > 100
```



“selection”

PName	Price	Manufacturer
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi




# Notation

Input Schema

Product(PName, Price, Category, Manufacturer)

```
SELECT PName, Price, Manufacturer
FROM   Product
WHERE  Price > 100
```



Answer(PName, Price, Manufacturer)

Output Schema



# Deletion Syntax

To delete rows from the table:

```
DELETE FROM <table name>  
WHERE <condition>;
```



# Deletion Example

```
DELETE FROM FoodCart  
WHERE food = 'hotdog';
```

FoodCart

date	food	sold
02/25/08	pizza	349
02/26/08	hotdog	500
02/26/08	pizza	70

Becomes

date	food	sold
02/25/08	pizza	349
02/26/08	pizza	70



Note: If the WHERE clause is omitted all rows of data are deleted from the table.

# Another Delete Example

```
DELETE FROM Student  
WHERE sNumber=6;
```



# Update Syntax

To update the content of the table:

*UPDATE* <table name>

*SET* <attr> = <value>

*WHERE* <selection condition>;



# Update Example

UPDATE FoodCart SET sold = 349

WHERE date = '02/25/08' AND food = 'pizza';

FoodCart

date	food	sold
02/25/08	pizza	350
02/26/08	hotdog	500
02/26/08	pizza	70



date	food	sold
02/25/08	pizza	349
02/26/08	hotdog	500
02/26/08	pizza	70



# Constraints

- Constraints are used to enforce the integrity of the data in a table by defining rules about values that can be stored in the columns of the table.
- Types of constraints include:
  1. NOT NULL – prevents null values from being stored in a column
  2. PRIMARY KEY
  3. FOREIGN KEY CONSTRAINT





# SQL NULL Values

- NULL values represent missing unknown data.
- By default, a table column can hold NULL values.
- If a column in a table is optional, we can insert a new record or update an existing record without adding a value to this column. This means that the field will be saved with a NULL value.



# SQL NULL Values

"Persons" table:

P_Id	LastName	FirstName	Address	City
1	Hansen	Ola		Sandnes
2	Svendson	Tove	Borgvn 23	Sandnes
3	Pettersen	Kari		Stavanger

- Suppose that the "Address" column in the "Persons" table is optional. This means that if we insert a record with no value for the "Address" column, the "Address" column will be saved with a NULL value.
- How can we test for NULL values?
- It is not possible to test for NULL values with comparison operators, such as =, <, or <>.
- We will have to use the **IS NULL** and **IS NOT NULL** operators instead.



# SQL IS NULL

- How do we select only the records with NULL values in the "Address" column?
- We will have to use the IS NULL operator:

```
SELECT LastName,FirstName,Address FROM Persons  
WHERE Address IS NULL;
```

The result-set will look like this:

LastName	FirstName	Address
Hansen	Ola	
Pettersen	Kari	



# PRIMARY KEY

- There are 2 ways to define a PK

1. Column-level constraint

Put the PRIMARY KEY keywords after the data type for the column

2. Table-level constraint

You can also define a constraint at the table level you can provide a name for the constraint

See examples on next slide



# Column-level Constraint

```
CREATE TABLE students
{
    student_id INT    PRIMARY KEY,
    student-name    VARCHAR(30) NOT NULL
};
```



# Table-level constraint

```
CREATE TABLE students
{
    student_id          INT,
    student-name        VARCHAR(30)    NOT NULL,
    CONSTRAINT student_pk PRIMARK KEY (student_id)
}
```

You can code PK constraint either way, both have the same effect



# FOREIGN KEY CONSTRAINT

- A *foreign key constraint* (reference constraint) requires values in one table to match values in another table. This defines the relationship between two tables and enforces referential integrity.



# A table with a column-level foreign key constraint

- To create a fk constraint at the column level, you code the **REFERENCES** keyword followed by the name of the related table and the name of the related column in parentheses.

```
CREATE TABLE invoices  
(  
    invoice_id    INT    PRIMARY KEY,  
    vendor_id     INT    REFERENCES vendors (vendor_id),  
    invoice_number VARCHAR(50) NOT NULL    UNIQUE  
);
```





# A table with a table-level foreign key constraint

```
CREATE TABLE invoices
(  
  invoice_id    INT      PRIMARY KEY,  
  vendor_id     INT      NOT NULL,  
  invoice_number VARCHAR(50) NOT NULL  UNIQUE,  
  CONSTRAINT invoices_fk_vendors  
    FOREIGN KEY (vendor_id)  
      REFERENCES vendors (vendor_id)  
);
```

Note if you try to Insert a row into invoices with a vendor\_id value that isn't matched by the vendor\_id column in the vendors table, you will get an error message



# Referential Integrity

- The system will display an error message if a constraint was violated
- If you try to delete rows that have related rows in another table , you will also see an error message
- In some cases you may want to automatically delete the related rows
  - use ON DELETE CASCADE



# Referential Integrity ctd

- Then when you delete a row from the pk table, the delete is cascaded to the related rows in the foreign table

CONSTRAINT invoices\_fk\_vendors

FOREIGN KEY (vendor\_id) REFERENCES vendors (vendor\_id)

ON DELETE CASCADE

- Use cascading deletes with caution ...



# Keys and Foreign Keys

Company

<u>CName</u>	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

Key

Product

<u>PName</u>	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Foreign  
key



# Altering the constraints of a table

- You may need to change the constraints of a table after you create it
- Use the ALTER TABLE statement

ALTER TABLE vendors

ADD PRIMARY KEY (vendor\_id)

ALTER TABLE vendors

DROP PRIMARY KEY



# SQL Queries:

## The LIKE operator

```
SELECT *  
FROM Products  
WHERE PName LIKE '%gizmo%'
```

- **LIKE** : pattern matching on strings
- may contain two special symbols:
  - % = any sequence of characters
  - \_ = any single character



# LIKE Examples

- Select firstnames where second letter = U
- Select \* from table where firstName like '\_u%';
- Select lastnames that end in S
- Select \* from table where lastName like '%s';



# Eliminating Duplicates

```
SELECT DISTINCT category  
FROM Product
```



Category
Gadgets
Photography
Household

Compared to:

```
SELECT category  
FROM Product
```



Category
Gadgets
Gadgets
Photography
Household





# SELECT DISTINCT Example

**"Persons" table:**

P_Id	LastName	FirstName	Address	City
1	Hansen	Ola	Timoteivn 10	Sandnes
2	Svendson	Tove	Borgvn 23	Sandnes
3	Pettersen	Kari	Storgt 20	Stavanger

SELECT **DISTINCT** City  
FROM Persons

City
Sandnes
Stavanger



# Another DISTINCT Example

name	city	director	assets
Branch_one	Jakarta	Bo Lee	80000
Clementi	Singapore	Ng Wee Hiong	3000000
F_branch	Johor Barhu	John	1500000
KL_branch	Kuala Lumpur	Yu Fei	1000000
Monas	Jakarta	Agus Arianto	4000000
S_branch	Johor Barhu	George	1200000

```
SELECT COUNT(DISTINCT city) as NumCities  
FROM branch
```

NumCities
4



# SQL ORDER BY Keyword

The ORDER BY keyword is used to sort the result-set by a specified column.

The ORDER BY keyword sort the records in ascending order by default.

If you want to sort the records in a descending order, you can use the DESC keyword.

## **SQL ORDER BY Syntax:**

```
SELECT column_name(s)  
  FROM table_name  
  ORDER BY column_name(s) ASC|DESC
```



# ORDER BY (Default ASC) Example

## "Persons" table:

P_Id	LastName	FirstName	Address	City
1	Hansen	Ola	Timoteivn 10	Sandnes
2	Svendson	Tove	Borgvn 23	Sandnes
3	Pettersen	Kari	Storgt 20	Stavanger

SELECT \* FROM Persons  
ORDER BY LastName

P_Id	LastName	FirstName	Address	City
1	Hansen	Ola	Timoteivn 10	Sandnes
4	Nilsen	Tom	Vingvn 23	Stavanger
3	Pettersen	Kari	Storgt 20	Stavanger
2	Svendson	Tove	Borgvn 23	Sandnes



# ORDER BY (DESC) Example

**"Persons" table:**

P_Id	LastName	FirstName	Address	City
1	Hansen	Ola	Timoteivn 10	Sandnes
2	Svendson	Tove	Borgvn 23	Sandnes
3	Pettersen	Kari	Storgt 20	Stavanger

SELECT \* FROM Persons  
ORDER BY LastName **DESC**

_Id	LastName	FirstName	Address	City
2	Svendson	Tove	Borgvn 23	Sandnes
3	Pettersen	Kari	Storgt 20	Stavanger
4	Nilsen	Tom	Vingvn 23	Stavanger
1	Hansen	Ola	Timoteivn 10	Sandnes



PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT DISTINCT category
FROM Product
ORDER BY category
```



Category
Gadgets
Household
Photography

```
SELECT Category
FROM Product
ORDER BY PName
```



Category
Gadgets
Household
Gadgets
Photography

```
SELECT DISTINCT category
FROM Product
ORDER BY PName
```



Category
Gadgets
Household
Photography



# SQL Alias

With SQL, an alias name can be given to a table.

## SQL Alias Syntax:

```
SELECT column_name(s)  
      FROM table_name  
      AS alias_name;
```



# SQL Scalar functions

- SQL scalar functions return a single value, based on the input value.
- Useful scalar functions:
  - UCASE() - Converts a field to upper case
  - LCASE() - Converts a field to lower case
  - LENGTH() - Returns the length of a text field
  - ROUND() - Rounds a numeric field to the number of decimals specified
  - NOW() - Returns the current system date and time





# SQL Aggregate Functions

- SQL has many built-in functions for performing calculations on data.
- SQL Aggregate Functions return a single value, calculated from values in a column. Useful aggregate functions:
  - AVG() - Returns the average value
  - COUNT() - Returns the number of rows
  - SUM() - Returns the sum



# Aggregation

```
SELECT avg(price)
FROM   Product
WHERE  maker="Toyota"
```

```
SELECT count(*)
FROM   Product
WHERE  year > 1995
```

SQL supports several aggregation operations:

sum, count, min, max, avg



# Aggregation: Count

COUNT applies to duplicates, unless otherwise stated:

```
SELECT Count(category)
FROM   Product
WHERE  year > 1995
```

same as Count(\*)

We probably want:

```
SELECT Count(DISTINCT category)
FROM   Product
WHERE  year > 1995
```



# Simple Aggregations

## Purchase

Product	Date	Price	Quantity
Bagel	10/21	1	20
Banana	10/3	0.5	10
Banana	10/10	1	10
Bagel	10/25	1.50	20

```
SELECT Sum(price * quantity)
FROM Purchase
WHERE product = 'bagel'
```



50 (= 20+30)



# SQL COUNT(column\_name) Example

## "Orders" table:

_Id	OrderDate	OrderPrice	Customer
1	2008/11/12	1000	Hansen
2	2008/10/23	1600	Nilsen
3	2008/09/02	700	Hansen
4	2008/09/03	300	Hansen
5	2008/08/30	2000	Jensen
6	2008/10/04	100	Nilsen

```
SELECT COUNT(Customer) AS CustomerNilsen  
FROM Orders  
WHERE Customer='Nilsen'
```

CustomerNilsen
----------------

2
---



# SQL COUNT(\*) Example

**"Orders" table:**

_Id	OrderDate	OrderPrice	Customer
1	2008/11/12	1000	Hansen
2	2008/10/23	1600	Nilsen
3	2008/09/02	700	Hansen
4	2008/09/03	300	Hansen
5	2008/08/30	2000	Jensen
6	2008/10/04	100	Nilsen

```
SELECT COUNT(*) AS NumberOfOrders  
FROM Orders
```

NumberOfOrders
6



# SQL COUNT(DISTINCT column\_name) Example

"Orders" table:

_Id	OrderDate	OrderPrice	Customer
1	2008/11/12	1000	Hansen
2	2008/10/23	1600	Nilsen
3	2008/09/02	700	Hansen
4	2008/09/03	300	Hansen
5	2008/08/30	2000	Jensen
6	2008/10/04	100	Nilsen

```
SELECT COUNT(DISTINCT Customer)
AS NumberOfCustomers
FROM Orders
```

NumberOfCustomers
-------------------

3
---



# SQL AVG() Function

The AVG() function returns the average value of a numeric column.

## **SQL AVG() Syntax:**

```
SELECT AVG(column_name)  
FROM table_name
```





# SQL AVG() Example

**"Orders" table:**

_Id	OrderDate	OrderPrice	Customer
1	2008/11/12	1000	Hansen
2	2008/10/23	1600	Nilsen
3	2008/09/02	700	Hansen
4	2008/09/03	300	Hansen
5	2008/08/30	2000	Jensen
6	2008/10/04	100	Nilsen

```
SELECT AVG(OrderPrice) AS OrderAverage  
FROM Orders
```

OrderAverage
--------------

950
-----



# SQL Cont'd

- Select (*Group By, Having Clause*)
- Now(), AND, OR
- Nested Subqueries
- More DML
  - Insert
  - Delete
  - Update
- Constraints



# How to group and summarize data

- The GROUP BY clause groups the rows of a result set based on one or more columns or expressions.
- Syntax

```
SELECT column_name, aggregate_function(column_name)
FROM table_name
WHERE column_name operator value
GROUP BY column_name;
```



# Example using details table

id	firstName	lastName	age	gender	position	department	rate	hours
1	Joe	Mullins	64	M	Lecturer	Engineering	63.08	12
2	Joan	Macgill	27	F	Researcher	Science	38.00	35
3	Jim	Mitchell	51	M	Researcher	Business	38.00	25
4	John	Magner	47	M	Lecturer	Humanities	63.08	16
5	Jean	Madden	45	F	Professor	Design	76.45	14
6	Jack	Minogue	61	M	Administrator	Hospitality	45.57	37
7	Josephine	Mahony	33	F	Head	Nursing	98.56	40
8	Juan	Mosley	56	M	Professor	Engineering	76.45	11
9	Jamie	Mullen	45	M	Researcher	Science	38.00	37

```
select department, count(*) AS "Number of Emp "  
from detailslab2.details  
group by department;
```



	department	Number of Emp
▶	Business	4
	Design	2
	Engineering	4
	Hospitality	3
	Humanities	1
	Nursing	2
	Science	4



# Group By with 2 columns

- If you include 2 or more columns or expressions in the GROUP BY clause, they will form a hierarchy where each column or expression is subordinate to the previous one.

```
select department,gender, count(*) AS "Number of Emp "  
from detailslab2.details  
group by department,| gender;
```



	department	gender	Number of Emp
	Business	F	2
	Business	M	2
	Design	F	2
	Engineering	F	1
	Engineering	M	3
	Hospitality	M	3
	Humanities	M	1
	Nursing	F	2
▶	Science	F	2
	Science	M	2



# Group By and Having Clause

- The GROUP BY clause determines how the selected rows are grouped, and the HAVING clause determines which groups are included in the final results



# Group By & Having Clause example

```
select department,gender, count(*) AS "Number of Emp "  
from detailslab2.details  
group by department,gender;
```



	department	gender	Number of Emp
	Business	F	2
	Business	M	2
	Design	F	2
	Engineering	F	1
	Engineering	M	3
	Hospitality	M	3
	Humanities	M	1
	Nursing	F	2
▶	Science	F	2
	Science	M	2

```
select department,gender, count(*) AS "Number of Emp "  
from detailslab2.details  
group by department, gender  
having count(*) >2;
```



	department	gender	Number of Emp
▶	Engineering	M	3
	Hospitality	M	3



# Group By & WITH ROLLUP

- You can use the WITH ROLLUP operator in the GROUP BY clause to add summary rows to the final result set.





# WITH ROLLUP example

select department, gender, count(\*) AS "Number of Emp"  
from details  
group by department, gender WITH ROLLUP;



	department	gender	Number of Emp
▶	Business	F	2
	Business	M	2
	Business	NULL	4
	Design	F	2
	Design	NULL	2
	Engineering	F	1
	Engineering	M	3
	Engineering	NULL	4
	Hospitality	M	3
	Hospitality	NULL	3
	Humanities	M	1
	Humanities	NULL	1
	Nursing	F	2
	Nursing	NULL	2
	Science	F	2
	Science	M	2
	Science	NULL	4
	NULL	NULL	20



# SQL NOW() Function

The NOW() function returns the current system date and time.

## SQL NOW() Syntax:

```
SELECT NOW()  
FROM table_name
```



# SQL NOW() Example

"Products" table:

Prod_Id	ProductName	Unit	UnitPrice
1	Jarlsberg	1000 g	10.45
2	Mascarpone	1000 g	32.56
3	Gorgonzola	1000 g	15.67

SELECT ProductName, UnitPrice, Now() as PerDate FROM Products

ProductName	UnitPrice	PerDate
Jarlsberg	10.45	10/7/2008 11:25:02 AM
Mascarpone	32.56	10/7/2008 11:25:02 AM
Gorgonzola	15.67	10/7/2008 11:25:02 AM



# SQL AND Operator

The AND operator is used to filter records based on more than one condition.

The AND operator displays a record if both the first condition and the second condition is true.



# AND Example

**"Persons" table:**

P_Id	LastName	FirstName	Address	City
1	Hansen	Ola	Timoteivn 10	Sandnes
2	Svendson	Tove	Borgvn 23	Sandnes
3	Pettersen	Kari	Storgt 20	Stavanger

SELECT \* FROM Persons  
WHERE FirstName='Tove'  
AND LastName='Svendson'

_Id	LastName	FirstName	Address	City
2	Svendson	Tove	Borgvn 23	Sandnes



# SQL OR Operator

The OR operator is used to filter records based on more than one condition.

The OR operator displays a record if either the first condition or the second condition is true



# OR Example

**"Persons" table:**

P_Id	LastName	FirstName	Address	City
1	Hansen	Ola	Timoteivn 10	Sandnes
2	Svendson	Tove	Borgvn 23	Sandnes
3	Pettersen	Kari	Storgt 20	Stavanger

```
SELECT * FROM Persons  
WHERE FirstName='Tove'  
OR FirstName='Ola'
```

P_Id	LastName	FirstName	Address	City
1	Hansen	Ola	Timoteivn 10	Sandnes
2	Svendson	Tove	Borgvn 23	Sandnes



# Nested Queries

- There may be scenarios where you need the output of one query to be used in another query. This is possible through nested queries.
  - For example. In the details lab, you may want to view staff who are younger than the average age.
  - Instead of using two separate queries as follows:
    - Select avg(age) from details; “which returns 41”
    - Followed by
    - Select \* from details where age < 41
- you can do the following:
- Select \* from details where age < (select avg(age) from details);